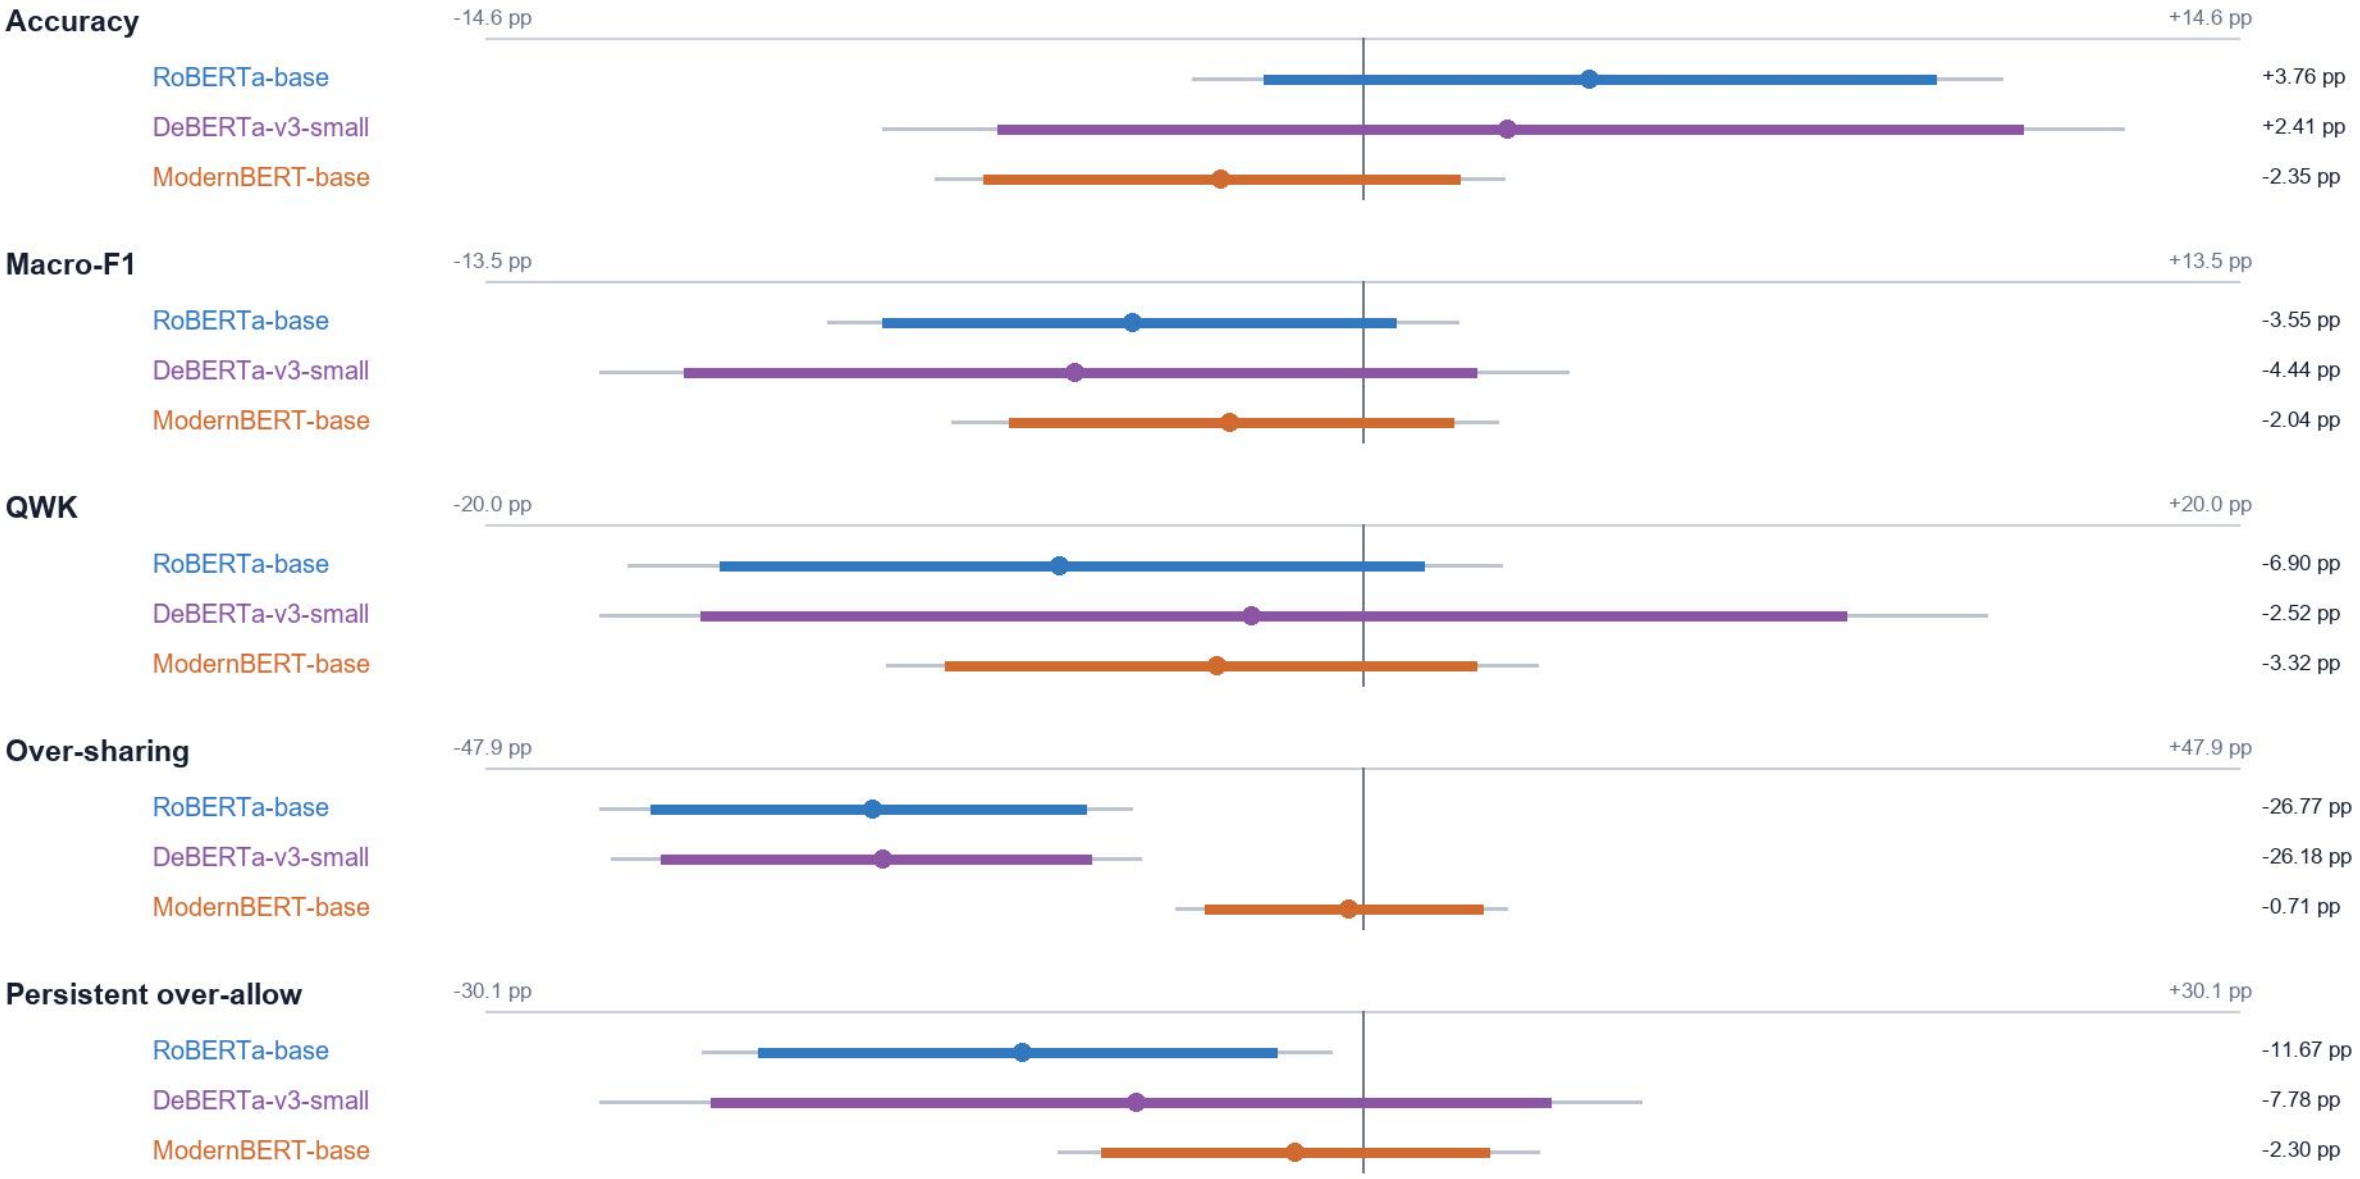
Semantic backbone vs tuned Flat LR: participant-paired differences

Values are direction-normalized: positive means the semantic model is better; risk-rate signs are reversed.



Thick line: marginal 95% CI. Thin gray line: Bonferroni 98.333% CI across 3 backbones per metric.

B=5,000; n=41 test participants; row-pooled metrics; intervals do not include training-seed or outer-split uncertainty.